



BFT in Lens of Blockchain

BFT IN LENS OF BLOCKCHAIN

Ted Yin^{1,2}, Dahlia Malkhi², Michael K. Reiter^{2,3}


Guy Golan Gueta² and Ittai Abraham²

¹Cornell University, ²VMware Research, ³UNC-Chapel Hill

About Me

- {Design,prov,build}ing **practical** distributed systems with **fundamental** (algorithmic) improvements
- Major work:
 - Avalanche Consensus (permission-less, extremely scalable)
 - HotStuff Consensus (permission-ed, elegant and drop-in replacement for PBFT/PBFT-like use cases)
- Two other on-going projects (Cornell, VMware)

BFT Consensus: Research In Our Eyes

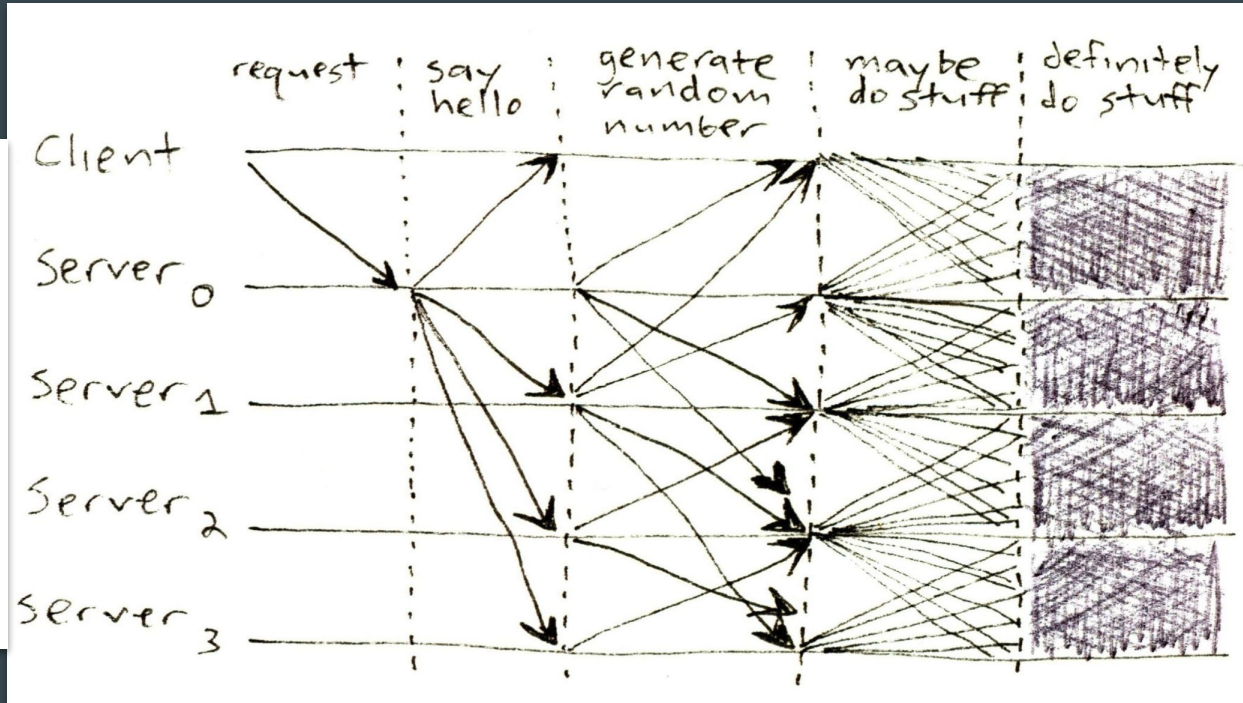



What we think we do

What others think we do

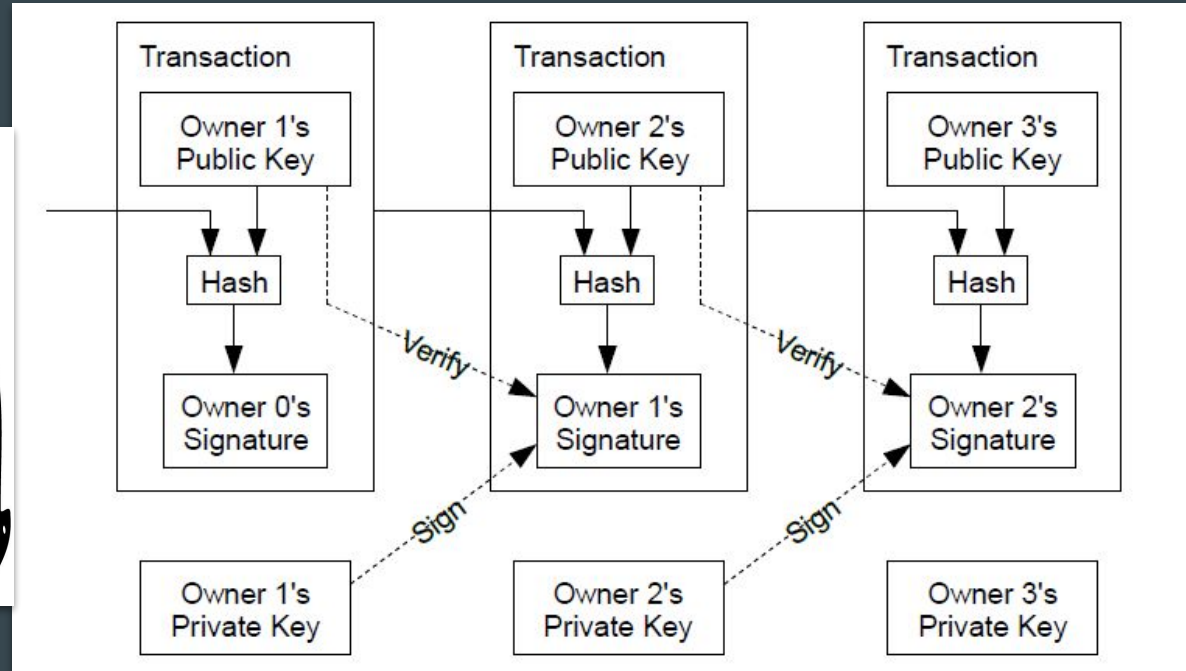



BFT Consensus: Why Another Protocol?



The Saddest Moment, Mickens 2013


BFT Consensus: Why Another Protocol?



Bitcoin: a Peer-to-Peer Electronic Cash System, Nakamoto 2008

BFT Consensus: Problem Definition

- N nodes replicate the same sequence of commands
- Consistent in asynchronous network (safety)
- During period of synchrony, it'd better progress (liveness)
- When the proposer (leader) is correct, it should be fast



Reducing the Complexity

“Communication Complexity”

Possibly the first
protocol with **linear cost**
during a view change

“Complexity”



Network Cost

“Protocol Complexity”

Protocol Spec

Conferences probably
don't care. But we do!

Protocol Complexity

Classical BFT

- **PoW-free**
- **Quorum invariants**
- From single-decree (1)
- (1) => Sequence numbers
- (1) => View numbers
- Hard to comprehend

Nakamoto's Consensus

- PoW based
- **Longest chain**
- **Naturally multi-decree (2)**
- (2) => **Block heights**
- (2) => **Views == Forks**
- **Easy to understand**

HotStuff: Protocol Framework & Simplicity

Framework

- Classical BFT variant (same/better guarantee)
- Bridges classical BFT and blockchain
- View change is everywhere, and nowhere
- Locking mechanism (reducing protocol state space)
- Decouples safety and liveness
- “Liveness gadget” could be RR, PoW based, etc.


Challenge: BFT consensus in 10 min

Ingredients to Make a 2-step HotStuff


- Protocol state variables
- Message types
- Voting rule
- Commit rule

Quorum Certificate

QC: Proof of the Existence of $2f+1$ (positive) Votes



Blockchain!

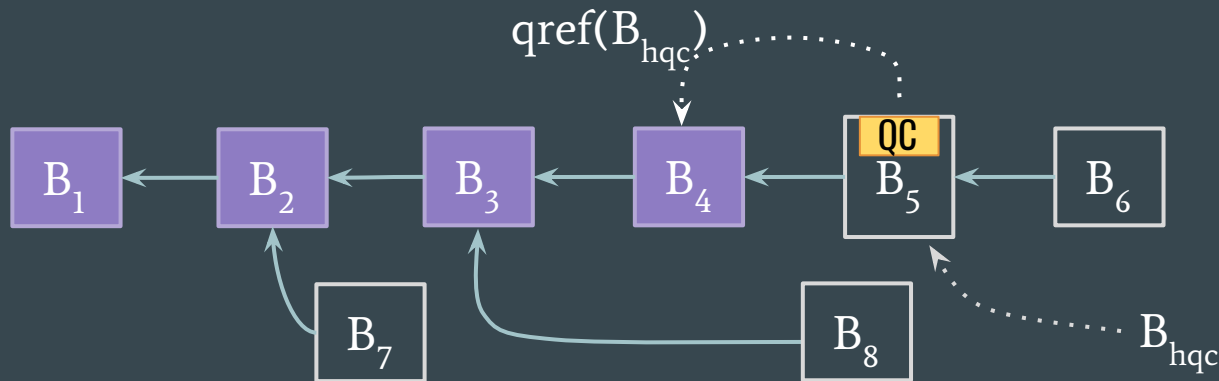


Branch Preference

B_{hqc} : block containing QC for the preferred block

“Preferred block” or $\text{qref}(B_{\text{hqc}})$: highest block receives a QC


Locking mechanism: a replica sticks to $\text{qref}(B_{\text{hqc}})$ unless...



Challenge: BFT consensus in 10 min

Protocol State Variables

- B_{hqc} = block containing a reference to the preferred branch
- B_{exec} = last committed block
- $vheight$ = height of the block last voted for



Challenge: BFT consensus in 10 min

Message Types

$\langle \text{propose}, v, B_{\text{new}}, B_{\text{hqc}} \rangle$


$\langle \text{vote}, \langle v, B_{\text{new}} \rangle_{\text{signed by } v}, B_{\text{hqc}} \rangle$

- Proposer broadcasts the propose message for block B_{new}
- Voters give back their opinions to the next proposer via votes
- **Only one type of messages for voting/view change, etc.**

Challenge: BFT consensus in 10 min

How to Vote?

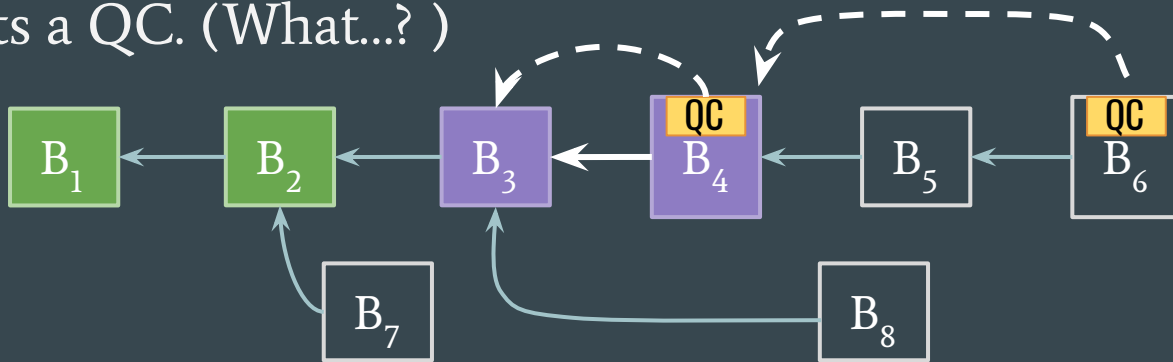
- Only vote positively for B_{new} if the following constraints hold:
 - $B_{\text{new}}.\text{height} > \text{vheight}$
 - B_{new} is on the same branch as $\text{qref}(B_{\text{hqc}})$



Challenge: BFT consensus in 10 min

When to Commit?

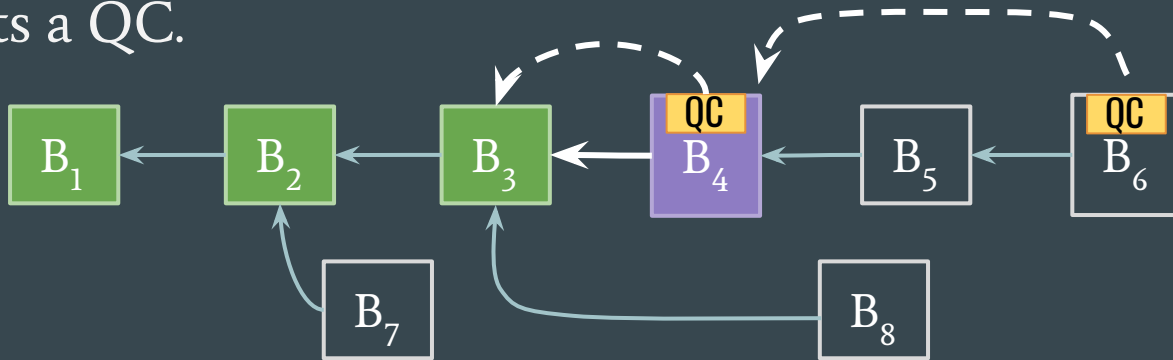
- Every block could contain a QC for some previous block
- Block B will be committed when a child having QC for B also gets a QC. (What...?)



Challenge: BFT consensus in 10 min

When to Commit?

- Every block could contain a QC for some previous block
- Block B will be committed when a child having QC for B also gets a QC.



HotStuff: Protocol in a Single Slide (2-step version)

Pseudo-code for replica u

```
1: // begin: rules specific to 2-step HotStuff in framework
2: function GETPREF() := QREF( $B_{\text{hqc}}$ )
3: function CHECKCOMMIT
4:   // check for a Commit 2-chain
5:    $B' := \text{QREF}(B_{\text{hqc}})$ 
6:    $B := \text{QREF}(B')$ 
7:   if  $B = B'.\text{parent}$  then
8:     ONCOMMIT( $B$ ); return true
9:   else return false
10: // end
11: // begin: generic HotStuff framework logic
12: procedure FINISHQC( $B$ )
13:    $\text{certs}[B] := \text{MAKEQC}(\text{ref} = B, \text{votes}[B])$ 
14: procedure ONCOMMIT( $B$ )
15:   if  $B_{\text{exec}}.\text{height} < B.\text{height}$  then
16:     ONCOMMIT( $B.\text{parent}$ )
17:     EXECUTE( $B.\text{cmd}$ )
18: procedure UPDATE( $B'_{\text{hqc}}$ )
19:   if  $\text{QREF}(B'_{\text{hqc}}).\text{height} > \text{QREF}(B_{\text{hqc}}).\text{height}$  then
20:      $B_{\text{hqc}} := B'_{\text{hqc}}$ 
21:   if CHECKCOMMIT then  $B_{\text{exec}} := B$ 
22: procedure ONRECEIVEPROPOSAL( $\langle \text{propose}, v, B_{\text{new}}, B'_{\text{hqc}} \rangle$ )
23:   UPDATE( $B'_{\text{hqc}}$ )
24:   if  $B_{\text{new}}.\text{height} > \text{vheight} \wedge \text{GETPREF}() \stackrel{*}{\leftarrow} B_{\text{new}}$  then
25:      $\text{vheight} := B_{\text{new}}.\text{height}$ 
26:      $\text{vote} := \langle \text{vote}, \langle u, B_{\text{new}} \rangle_{\sigma_u}, B_{\text{hqc}} \rangle$ 
27:     SEND(NEXTPROPOSER( $v$ ),  $\text{vote}$ )
28: procedure ONRECEIVEVOTE( $\langle \text{vote}, \langle v, B_{\text{new}} \rangle_{\sigma_v}, B'_{\text{hqc}} \rangle$ )
29:   UPDATE( $B'_{\text{hqc}}$ )
30:   if  $\exists \langle v, B_{\text{new}} \rangle_{\sigma'_v} \in \text{votes}[B_{\text{new}}]$  then return
31:   // collect votes for  $B_{\text{new}}$ 
32:    $\text{votes}[B_{\text{new}}] := \text{votes}[B_{\text{new}}] \cup \{ \langle v, B_{\text{new}} \rangle_{\sigma_v} \}$ 
33:   if  $|\text{votes}[B_{\text{new}}]| \geq 2f + 1$  then FINISHQC( $B_{\text{new}}$ )
34: procedure ONPROPOSE( $B_{\text{tail}}, \text{qc}, \text{cmd}$ )
35:    $B_{\text{new}} := \text{MAKEBLOCK}(\text{parent} = B_{\text{tail}},$ 
36:      $\text{height} = B_{\text{tail}}.\text{height} + 1,$ 
37:      $\text{qc} = \text{qc}, \text{cmd} = \text{cmd})$ 
38:   // send to all replicas, including  $u$  itself
39:   BROADCAST( $\langle \text{propose}, u, B_{\text{new}}, B_{\text{hqc}} \rangle$ )
40: // end
```

HotStuff: Protocol in a Single Slide (3-step version)

```
1: // begin: rules specific to 3-step HotStuff in framework
2: function GETPREF() := QREF(QREF( $B_{\text{hqc}}$ ))
3: function CHECKCOMMIT
4:   // check for a Commit 3-chain
5:    $B'' := \text{QREF}(B_{\text{hqc}})$ 
6:    $B' := \text{QREF}(B'')$ 
7:    $B := \text{QREF}(B')$ 
8:   if ( $B = B'.\text{parent}$ )  $\wedge$  ( $B' = B''.\text{parent}$ ) then
9:     ONCOMMIT( $B$ ); return true
10:  else return false
11: // end
```

HotStuff vs. State of the Art Performance




Figure 9: Throughput vs. number of nodes with payload size 0/0 and 1024/1024.




Figure 10: Latency vs. number of nodes with payload size 0/0 and 1024/1024.

HotStuff vs. State of the Art Performance

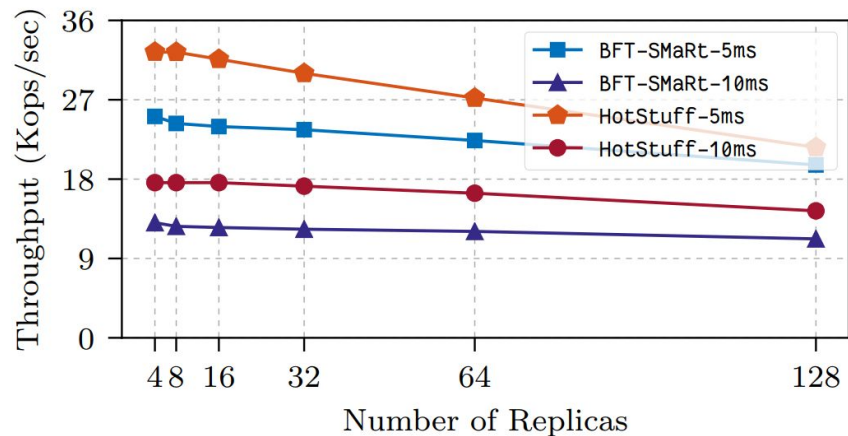


Figure 11: Throughput vs. number of nodes with inter-replica latency 5ms and 10ms.


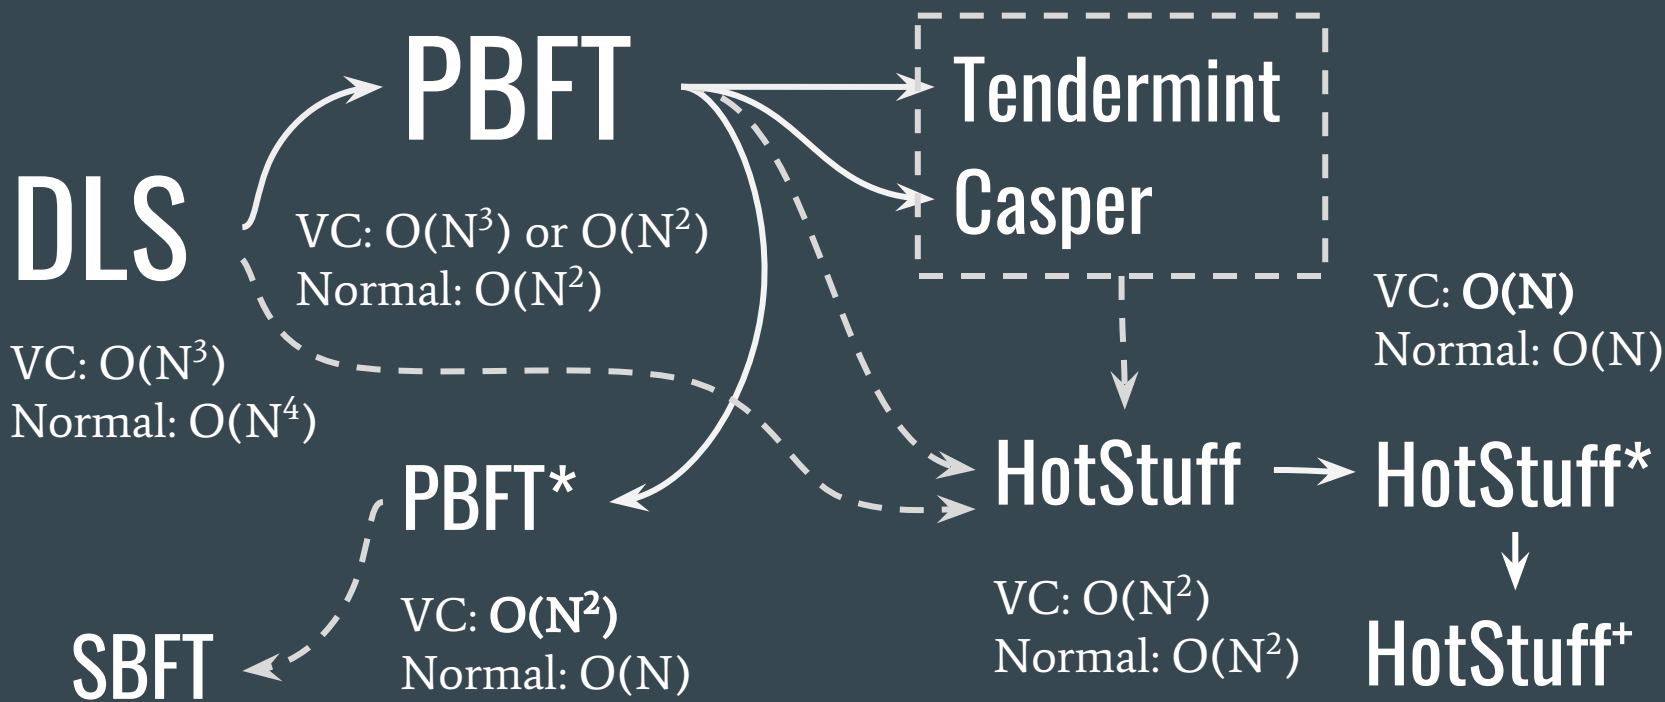



Figure 12: Latency vs. number of nodes with inter-replica latency 5ms \pm 0.5ms or 10ms \pm 1.0ms.

BFT Solutions & Communication Complexity

VC = View Change



That'd be all.
Special Thanks
VMware Research Group



<https://arxiv.org/abs/1803.05069>

Open-sourced code coming soon